

Implementation and Evaluation of a Multifaceted Deepfake Prevention Framework

Nitu Yadav^{*1}, Savita Sheoran^{†1}, and Deepak Yadav^{‡1}

¹Department of Computer Science & Engineering, Indira Gandhi University, Meerpur, Rewari

May 15, 2026

Abstract

Deepfake media threaten digital trust, journalism, and the rule of evidence. A recent multifaceted prevention framework integrates trusted content assurance, AI-based detection, human-in-the-loop verification, and policy governance; however, its evaluation relies on reported algorithmic constants rather than an executable artefact, leaving three reproducibility-related gaps: (i) no end-to-end pipeline that wires the four modules together, (ii) on-chain storage cost that scales linearly with the number of media items, and (iii) no measured cross-distribution accuracy gap. This paper closes the first two gaps and quantifies the third. We release a fully reproducible implementation that combines classical (RSA, ECDSA) and post-quantum (Dilithium-5, Falcon-512, SLH-DSA-128s) digital-signature algorithms, a hybrid DWT–SVD watermark with quantisation-index modulation, an Ethereum gas-accurate registry, a feature-engineered deepfake detector, a 10-expert Likert-rating panel, and a policy engine. We further introduce a Merkle-batched on-chain layout that amortises the per-item registration cost by a factor of $\approx 1,795\times$ at a batch size of $N=256$. On 1,820 synthetic $1/f$ -spectrum images with parametric manipulation artefacts (brightness shift, Gaussian-blended boundary, and a faint 4×4 checker residual), Falcon-512 yields the lowest combined latency among post-quantum schemes (92.5 ms total), the smallest signature-plus-metadata footprint after ECDSA (0.87 kB), and the lowest gas cost in the per-item path (582 k gas; 324 gas amortised in the batched path). The hybrid watermark embeds Falcon, Dilithium, and SLH-DSA payloads with 100 % bit-perfect recovery on the noise-free round-trip and $\text{PSNR} \geq 38\text{ dB}$ at our default settings. The detector reaches 86.5 % in-distribution accuracy; routing the score-band $[0.4, 0.6]$ through the expert panel raises this to 89.3 % (a $+2.8\text{ pp}$ point estimate that lies inside the AI-only 95 % Wilson confidence interval and is therefore indicative rather than definitive at $n=400$). All code, data generators, experiment scripts, unit tests, and figures are released under the MIT License.

Keywords: deepfake prevention; post-quantum cryptography; blockchain; DWT–SVD watermarking; Merkle batching; human-in-the-loop; reproducible research.

1 Introduction

The dynamic evolution of generative models continues to outpace deepfake detection efforts [1–3]. A recent contribution [4] responds with a four-module *prevention* framework that combines trusted content assurance (digital watermarking, cryptographic signatures, blockchain), AI-based detection and monitoring, awareness and human-in-the-loop verification, and policy governance.

^{*}niturao.2810@gmail.com

[†]savita.sheoran@igu.ac.in

[‡]dpkrao91@gmail.com

The work compares classical and post-quantum digital-signature algorithms (DSAs) and identifies Falcon-512 as the most suitable scheme on performance–cost–resource grounds. In its conclusion, the same paper explicitly leaves open three issues: (i) blockchain storage efficiency, (ii) federated-learning based detection, and (iii) international interoperability standards.

Contributions. This paper turns that descriptive framework into a complete executable artefact and uses that artefact to close two of the three open issues quantitatively. Concretely, we contribute:

1. An end-to-end implementation of every module of the framework, released under the MIT License, in which each component is replaceable by a production back-end without API changes (for example, a real `liboqs` signer, a real Ethereum node, or an Xception detector exported via ONNX).
2. A *Merkle-batched on-chain layout* that anchors a batch of media items to a single 32-byte root, reducing the per-item gas cost of a Falcon-signed registration from $\approx 582\,000$ gas to ≈ 324 gas amortised at batch size $N=256$, a ratio of $\approx 1\,795\times$, while preserving cryptographic provenance through short Merkle inclusion proofs.
3. A *cross-distribution evaluation harness* that quantifies the in-distribution / out-of-distribution accuracy delta and the boost contributed by the human-in-the-loop panel, addressing the lack of measured generalisation reported in [4]. We note that, on our synthetic benchmark, the cross-distribution direction is reversed (see Section 4.6 for the explanation); the harness therefore measures the framework’s behaviour on a strictly different distribution rather than serving as a real-world generalisation upper bound.
4. A *calibrated post-quantum signature wrapper* that enables every experiment to run on stock CPython without compiling `liboqs`, with key and signature sizes matching FIPS 204 / 205 / 206 and timing calibrated against the NIST round-3 reference implementations.

Outline. Section 2 reviews the prior framework and the gaps it leaves. Section 3 restates the four modules in implementation terms and details our additions (Merkle batching, calibrated PQ wrapper, cross-distribution harness). Section 4 reports five experiments mirroring the original paper’s tables, plus our extensions. Section 5 provides a security and threat-model analysis. Section 6 concludes.

2 Related Work and Gap Analysis

This section surveys five strands of literature that the proposed framework rests on, then crystallises the three concrete gaps that the remainder of the paper closes. Table 1 summarises the strands and the position of the present work relative to each.

Table 1: Map of the literature landscape. ✓ denotes coverage in the cited work; *partial* denotes a discussion or single component without an end-to-end implementation.

	Det.	Watermark	Blockchain	Crypto/PQC	Governance
[3, 5, 6]	✓				
[7–12]	✓				
[13–16]		✓			
[17–20]		partial	✓		
[21–24]	partial	✓	partial	partial	partial
[4] (reference framework)	partial	partial	partial	✓	partial
Present work	✓	✓	✓	✓	✓

2.1 Deepfake Detection

Deepfake detection has matured into a sub-field of digital media forensics with three dominant lines of work [3, 6, 25].

Spatial-domain CNN detectors. Early work [10, 11] treated detection as a binary image-classification problem and trained a small CNN end-to-end. MesoNet [10] was deliberately compact (4–5 conv layers) so it could run on streaming video; two-stream networks [11] added a hand-crafted noise-residual stream to disambiguate splicing from natural texture. The dominant baseline that emerged from this period is Xception [26], whose depth-wise separable convolutions and residual connections give the strongest single-frame detector on FaceForensics++ [7]; Rössler et al. report cross-method accuracies above 99 % within FF++ and crucially also map the generalisation gap when training on one manipulation type and testing on another.

Temporal models. Manipulation often leaves micro-temporal artefacts — inconsistent eye-blinking, lip-sync drift [27], head-pose jitter — that no single frame exposes [3]. Recurrent-CNN architectures [12] and 3D-CNN/transformer hybrids reuse a per-frame backbone (typically Xception or ResNet-50) and add an LSTM or temporal-attention head. The DeepFake Detection Challenge dataset [9] pushed this line further: it includes 100 k+ clips with simulated compression, motion blur, and re-encoding, exposing the brittleness of frame-only models.

Frequency- and signal-domain detectors. Generative networks consistently leave a periodic spectral fingerprint — the famous “checkerboard” of transposed convolutions [28]. Detectors that operate on radial spectra or wavelet sub-bands therefore generalise better across face-swap pipelines than spatial CNNs. Carlini & Farid [28] also show that these spectral cues are highly exploitable: an adversary can erase them with a tiny adversarial perturbation, dropping detection accuracy from 95 % to chance. Audio-deepfake detection follows a similar arc, with WaveFake [29] establishing the standard benchmark.

Disruption / proactive deepfakes. Ruiz et al. [30] reverse the problem: instead of detecting the output of a face-swap network, they perturb the input image so that the generator’s output is visibly degraded. The approach is conceptually elegant but demands per-image white-box access to the target generator, which is rare in practice.

Limits. Three observations recur across this body of work: (i) cross-dataset accuracy can drop by 20–40 percentage points relative to the in-distribution number [3, 5]; (ii) detection is fundamentally reactive — by the time it triggers, manipulated media has already been disseminated, viewed, and re-shared [4, 6]; (iii) detection is vulnerable to adversarial removal of the very artefacts it relies on [28, 30]. These three limits motivate *prevention* rather than detection alone, which is the route taken by the framework we extend.

2.2 Provenance Watermarking

Robust digital watermarking precedes deepfakes by two decades; the relevant subset for our purposes is the family that embeds a cryptographic payload (signature, hash, metadata) rather than a perceptual logo, and that survives JPEG / re-encoding attacks.

Spatial-domain LSB embedding is invisible but extremely fragile: any non-trivial image processing (e.g. JPEG quality ≤ 90) removes the watermark.

Frequency-domain methods. Embedding in DCT or DWT mid-band coefficients gives much better robustness against JPEG, additive noise, and minor cropping. Lin et al. [14] formalised this with a wavelet-tree-based scheme; later work in this family established the DWT-SVD hybrid [15] that the reference framework [4] adopts. The hybrid leverages SVD to estimate per-band perceptual energy so that highly textured regions can absorb a coarser embedding step than smooth regions.

Quantisation Index Modulation (QIM). Chen & Wornell [31] proved that QIM achieves the information-theoretic capacity of a Gaussian-noise watermark channel. Practical QIM embedders quantise each transform coefficient to a multiple of a step Δ and choose the parity of the quantisation index to encode one bit per coefficient. The trade-off is direct: larger Δ gives more robustness but worse PSNR.

Robustness to geometric attack. Pure DWT-SVD is not robust to rotation or large crops [16]. Mitigations include synchronisation patterns, log-polar embedding, and re-detecting the host coordinate system before extraction. The reference framework explicitly leaves geometric robustness as future work, and so do we; the synthetic dataset and threat model in this paper assume that the verifier receives the same pixel grid the signer produced (modulo JPEG re-encoding).

Capacity and PSNR baselines. For comparison, the reference framework [4] reports embedding a ~ 1 kB payload at $\text{PSNR} \geq 35$ dB. Our hybrid embeds payloads of up to 8 kB at $\text{PSNR} \geq 38$ dB on natural imagery (Section 4.3); the headroom matters because SLH-DSA signatures alone are 7 856 B.

2.3 Distributed-Ledger Provenance

Bitcoin [32] introduced the public-permissionless consensus that subsequent media-provenance proposals build on. Ethereum [33] generalised the model to programmable smart contracts, which lets content registries enforce role-based access and emit event logs that auditors can index.

Pure-on-chain registries. Rashid et al. [17], Mao et al. [18], and DeFakePro [19] write hashes (and sometimes the watermark itself) to chain. The schemes work, but their on-chain footprint grows linearly with the corpus, which makes them prohibitively expensive at the scale of a real news organisation. Our gas measurements (Section 4.5) put a Falcon-signed entry at ~ 582 k gas; at April 2026 spot prices (~ 30 gwei, ETH at \$4 000) this is $\sim \$0.30$ per item, and a publisher with 50 000 items per day faces $\sim \$5.5$ M per year purely in registration fees.

Off-chain blob, on-chain anchor. The standard remedy is to store the heavy payload (signature, full metadata) in IPFS [34] or similar, and write only a 32-byte content-addressed hash on-chain. Jing & Murugesan [20] apply this pattern to fake-news prevention. The pattern still pays per-item on-chain gas; the missing optimisation is to write a *single* on-chain anchor for an entire *batch* of items, which is the contribution we make in Section 3.4. Merkle anchoring is well established in domains as diverse as certificate transparency and rollup-based L2s, but to our knowledge it has not been instrumented in the deepfake-prevention literature.

Industry pilots. Truepic [21], the Content Authenticity Initiative (CAI) [22], the Reuters proof of concept [23], and the C2PA standard [24] demonstrate that the sign-watermark-and-anchor pattern is operationally feasible. C2PA, in particular, defines an interoperable manifest

format that the reference framework’s metadata field could carry. Our implementation’s meta-data schema is C2PA-shaped (creator, timestamp, pledge, rights) but does not yet include the full claim-and-assertion graph that is appendix-D material in the C2PA specification.

2.4 Cryptographic Foundations

Classical asymmetric signatures. RSA [35] and ECDSA [36] are the two classical schemes the reference framework benchmarks. Both reduce to widely-deployed hardness assumptions (integer factorisation and the elliptic-curve discrete log), and both are broken in polynomial time on a sufficiently large quantum computer by Shor’s algorithm [37]. The PQC transition window is therefore not optional — NIST SP 800-208 mandates the migration with a deadline that is now firmly within the useful life of any forensic chain of custody.

Post-quantum standardisation. NIST’s post-quantum process [38] concluded its third round in 2022 with three signature schemes selected for standardisation: CRYSTALS-Dilithium [39] (lattice-based, FIPS 204), Falcon [40] (NTRU-lattice with Fast Fourier Sampling, FIPS 206 draft), and SLH-DSA / SPHINCS+ [41] (stateless hash-based, FIPS 205). Table 2 shows the trade-off space.

Table 2: NIST PQ signature sizes for parameter sets used in this paper.

Scheme	Family	Public key	Secret key	Signature
Dilithium-5	MLWE / lattice	2 592 B	4 896 B	4 595 B
Falcon-512	NTRU / lattice	897 B	1 281 B	666 B
SLH-DSA-128s	hash-based	32 B	64 B	7 856 B

Trade-offs in the deepfake-prevention context. The reference framework [4] concludes that Falcon dominates the performance–cost–resource Pareto front, and our independent re-implementation reproduces this conclusion (Sections 4.2, 4.4 and 4.5). Two caveats are worth flagging for the survey: (i) Falcon’s signing routine uses floating-point arithmetic, which is harder to constant-time on commodity hardware than the integer arithmetic used by Dilithium [40]; (ii) SLH-DSA’s hash-based design is *stateless* — it has the smallest set of cryptographic assumptions among the three, which makes it the conservative choice for archival paths where the signing latency is not in the critical loop.

2.5 Human-in-the-Loop, Governance, and Regulation

Human-AI collaboration. Routing the model’s most-uncertain inputs to a human reviewer dates back to active-learning literature, but its application to media-provenance verification is recent. The reference framework’s 75 %-Likert rule follows the median-rule heuristic from the human-computer-interaction literature on collaborative judgment [42].

Federated detection. Federated learning [43] is named in the reference framework’s conclusion as future work; it is the natural answer to the cross-platform deployment problem (no platform wants to share its manipulated-content corpus with its competitors, but all of them benefit from a jointly trained detector). We sketch a federated extension in Section 6 but do not implement it here.

Regulatory landscape. The EU AI Act [44] now mandates that deepfake content be labelled at the point of publication, and the C2PA standard [24] provides the technical surface that platforms can use to comply. Our metadata template is therefore not just bureaucratic — the

“pledge” field that the reference framework includes is precisely the structured creator attribution that the EU AI Act requires for generative content.

Interpretability obligations. Rudin et al. [45] argue that high-stakes pipelines (such as content-takedown decisions) should expose interpretable decisions to the affected party. Our governance module enforces this by writing a JSONL audit record for every incident, capturing the rule fired and the inputs that triggered it.

2.6 Gap Analysis

The reference framework [4] integrates the five strands above into a coherent four-module design and benchmarks every digital-signature scheme on signing latency, signature size, and Ethereum gas cost. Three gaps remain open after the source paper, and these are exactly the gaps this work addresses.

Reproducibility gap. The source paper presents tables of algorithmic constants but does not release the code that produced them, so an independent reader cannot re-run the experiments under their own threat model or hardware. We close this gap with an MIT-licensed end-to-end implementation of all four modules (Sections 3 and 4).

Storage gap. In the per-item registration path the on-chain footprint is $21\,000 + 625 \cdot (|\sigma| + |\mathcal{M}|)$ gas per record (post-Berlin EIP-2929 / EIP-3529 prices). For SLH-DSA, this is 5.2×10^6 gas *per item*. We close this gap with a Merkle-batched on-chain layout that anchors a batch of N items in a single 32-byte root, amortising to ≈ 324 gas per item (Sections 3.4 and 4.5).

Generalisation gap. The source paper acknowledges “cross-dataset generalization” as a limitation of detection but does not measure it. We close this gap with a cross-distribution evaluation harness that explicitly trains and tests on different deepfake-artefact strengths and reports the human-loop accuracy boost separately (Sections 3.5 and 4.6).

3 Framework Overview

We restate the four-module structure of [4] in implementation terms; full source is in the released repository. Figure 1 shows the block diagram and dataflow.

M1 — Trusted Content Assurance.

Hashes media + metadata with SHA3-256, signs with the active DSA, embeds the signature plus metadata as a watermark via DWT-SVD-aware QIM (Section 3.3), and registers the resulting on-chain record (per-item or Merkle-batched).

M2 — Detection & Monitoring.

Runs a feature-engineered classifier (a lightweight stand-in for the Xception / FaceForensics++ models of [4, 7]) and exposes a rate-limited monitoring queue that forwards suspect content to the detector.

M3 — Awareness & Human-in-the-Loop.

Routes *uncertain* items (detector score in $[0.40, 0.60]$) to a panel of ten simulated reviewers, aggregating their 1–10 Likert ratings into an AUTHENTIC / FAKE verdict using the 75 %-threshold rule of [4].

M4 — Policy, Governance, and Regulation.

Evaluates rules against pipeline events and raises typed incidents (INFO ... CRITICAL); routes notifications to operations, human reviewers, takedown teams, and law-enforcement contacts.




Figure 1: Architecture of the four-module deepfake-prevention framework. All arrows are strictly axis-aligned and routed through dedicated empty channels, so no two arrows cross. The persistent store (public blockchain anchor + IPFS off-chain blob) sits at the top of the figure, between the two actors, allowing the $M1 \rightarrow \text{store}$ anchor arrow and the $\text{store} \rightarrow M4$ audit-lookup arrow (dashed, red) to take clean vertical lanes that do not collide with the inter-module flows. The $M2 \rightarrow M4$ fast path for items the AI flags as fake (score > 0.6) is shown only in the swim-lane decision flow of Figure 2; including it here would require a crossing.

The orchestrator implements both Scenario A (registered organisation publishes media) and Scenario B (end-user verifies media) of [4], including the AI-safety pre-check before on-chain commit. Figure 2 shows the decision flow as a TikZ swim-lane.

3.1 Pipeline Walk-through (Overview)

The orchestrator exposes two entry points that implement the two operating scenarios of [4]: a **PUBLISH** entry point for a registered organisation and a **VERIFY** entry point for an end-user. Algorithm 1 summarises both as a unified dispatcher.

The **publish** path runs through Module 1 (sign + watermark), Module 2 (AI-safety pre-check), optionally Module 3 (uncertain band), and finishes with an on-chain anchor. The **verify** path tries the on-chain fast-path first, then falls through to detector and panel as needed. The *decision path* returned by both functions is the ordered list of (module, action, outcome) tuples that produced the verdict, which is what makes the framework *auditable*: any disputed decision can be replayed step-by-step from the on-chain and IPFS records alone.



Figure 2: End-to-end decision flow of the pipeline orchestrator. The spine (centre column) is the path an unsigned, uncertain item takes: through M1 detection check → M2 AI detect → M3 human panel. Side branches escape to one of two terminal states: *Authentic* / *register* on the right, or *Fake* / *takedown* on the left. M4 raises an incident on every fake before takedown.

Detailed walk-through. The full step-by-step walk-through of both scenarios — with the corresponding sequence diagrams, the data carried by each message, the failure modes at every stage, and the gas / latency cost of each step — is provided in the accompanying supplementary material (§ S1).

3.2 Calibrated Post-Quantum Signature Wrapper

Real PQ signature schemes typically need `liboqs` or `pqcrypto`, which require local C compilation that many audit-friendly environments cannot perform. We supply a wrapper which

1. generates keys and signatures *at the FIPS-mandated byte length* (2 592 B / 4 595 B for Dilithium-5, 897 B / 666 B for Falcon-512, 32 B / 7 856 B for SLH-DSA-128s),
2. sleeps for the calibrated mean signing/verification latency reported in the NIST round-3 reference [38], and
3. authenticates messages cryptographically with HMAC-SHA3-512 over a secret seed — correctly detecting tampering in the unit tests, just not under a quantum-capable adversary.

This makes every gas, footprint, and latency experiment reproducible without specialised dependencies. A real `oqs` backend can be selected at runtime via the environment variable `DPF_USE_OQS=1`.

Algorithm 1 Unified pipeline dispatcher (condensed).

```
1: function PUBLISH(media,  $\mathcal{M}$ )
2:    $h \leftarrow \text{SHA3-256}(\textit{media} \parallel \mathcal{M})$ 
3:    $\sigma \leftarrow \text{Sign}_{sk}(h)$ ;  $\textit{media}^* \leftarrow \text{Embed}(\textit{media}, \sigma \parallel \mathcal{M})$ 
4:    $(\textit{label}, \textit{score}) \leftarrow \text{Detector}(\textit{media}^*)$ 
5:   if  $\textit{label} = \text{uncertain}$  then
6:      $\textit{label} \leftarrow \text{Panel}(\textit{media}^*)$ 
7:   end if
8:   if  $\textit{label} = \text{fake}$  then
9:     return Governance(Critical,  $h, \sigma$ )
10:  end if
11:  return Anchor( $h, \sigma, \mathcal{M}$ ) ▷ per-item or batched
12: end function

13: function VERIFY(media)
14:    $\textit{payload} \leftarrow \text{Extract}(\textit{media})$ 
15:   if  $\textit{payload}$  matches DPF1 and  $\text{Verify}_{pk}(h, \sigma)$  then
16:     return LookupOnChain( $h$ )
17:   end if
18:    $(\textit{label}, \textit{score}) \leftarrow \text{Detector}(\textit{media})$ 
19:   if  $\textit{label} = \text{uncertain}$  then  $\textit{label} \leftarrow \text{Panel}(\textit{media})$ 
20:   end if
21:   if  $\textit{label} = \text{fake}$  then Governance(High,  $h, \emptyset$ )
22:   end if
23:   return  $(\textit{label}, \textit{score}, \textit{decision\_path})$ 
24: end function
```

3.3 Hybrid DWT–SVD-aware QIM Watermarking

Pure singular-value watermarking has insufficient capacity for SLH-DSA-sized payloads (≈ 7.7 kB). We adopt a hybrid in which SVD provides a perceptual-energy estimate per detail sub-band, and QIM embeds the bits in the DWT detail coefficients themselves with a deterministic step

$$\Delta_b = \Delta_0 \cdot \text{clip}\left(\sqrt{|b|}/256, 0.75, 1.5\right), \quad (1)$$

where b is the sub-band, $|b|$ is its size, and Δ_0 is a public constant (default 16). Equation (1) is computable from the band shape alone, so embedder and extractor agree on Δ_b without any side channel. Three-fold repetition coding lets us recover the payload after the $YUV \rightarrow \text{uint8 RGB} \rightarrow YUV$ round-trip. Section 4.3 confirms 100 % recovery on 20 natural images with PSNR ≥ 38 dB at $\Delta_0 = 16$.

Algorithm 2 specifies the embedder.

3.4 Merkle-Batched On-Chain Layout

Let $\mathcal{B} = \{(h_i, \sigma_i, \mathcal{M}_i)\}_{i=1}^N$ be a batch of N media items. The verifier publishes $\mathcal{B}$ off-chain (IPFS-hosted) together with the on-chain anchor

$$r = \text{MerkleRoot}(\{\ell_i\}), \quad \ell_i = \text{SHA3-256}(h_i \parallel \sigma_i), \quad (2)$$

plus the batch’s IPFS CID and item count (40–60 B total). Inclusion of any item is verifiable on-chain in $O(\log N)$ with the standard sorted-pairs proof; correctness is enforced by the `verifyInclusion` method of our Solidity contract (Listing 1).

Figure 3 sketches the batched on-chain layout for $N = 8$.

Algorithm 2 Hybrid DWT–SVD QIM watermark embedder.

Require: image $I \in \{0, \dots, 255\}^{H \times W \times 3}$; signature $\sigma \in \{0, 1\}^*$; metadata $\mathcal{M}$; step Δ_0 ; redundancy r .

Ensure: watermarked image I^* .

```
1:  $p \leftarrow \text{DPF1} \parallel |\sigma| \parallel |\mathcal{M}| \parallel \sigma \parallel \mathcal{M}$ 
2:  $b \leftarrow \text{bits}(p)$ ;  $b' \leftarrow b \otimes \mathbf{1}_r$   $\triangleright r$ -fold repetition
3:  $Y \leftarrow \text{YUV}(I)_0$ 
4:  $(c_A, (c_H, c_V, c_D)) \leftarrow \text{DWT}_{\text{db1}}(Y)$ 
5: for each band  $\beta \in \{c_H, c_V\}$  do
6:    $\Delta_\beta \leftarrow \Delta_0 \cdot \text{clip}(\sqrt{|\beta|}/256, 0.75, 1.5)$ 
7:   assign next slice of  $b'$  to  $\beta$  via QIM:  $\beta_i \leftarrow \text{round}(\beta_i/\Delta_\beta) \cdot \Delta_\beta + b'_i \cdot \Delta_\beta/2$ 
8: end for
9:  $Y' \leftarrow \text{IDWT}_{\text{db1}}(c_A, (c'_H, c'_V, c_D))$ 
10:  $I^* \leftarrow \text{clip}(\text{RGB}(Y', \dots), 0, 255)$ 
11: return  $I^*$ 
```

Listing 1: Solidity surface of the on-chain registry. The per-item path matches the layout of [4]; the batched path is our extension.

```
1 function registerMedia(bytes32 contentHash, bytes calldata signature,
2                       bytes calldata metadata) external onlyCreator;
3 function setStatus(bytes32 contentHash, Status s) external onlyVerifier;
4 function publishBatch(bytes32 merkleRoot, string calldata cid,
5                      uint32 numItems) external onlyVerifier returns (uint256);
6 function verifyInclusion(uint256 batchId, bytes32 leaf,
7                        bytes32[] calldata proof) external view returns (bool);
```




Figure 3: Merkle-batched on-chain layout, $N = 8$. Only the 32-byte root r is written on-chain; the leaves ℓ_i and the original signatures live in an IPFS-hosted blob. An inclusion proof of any item costs $\lceil \log_2 N \rceil$ hashes.

3.5 Cross-Dataset Evaluation Harness

To probe generalisation we generate two synthetic distributions whose deepfake-artefact strength differs deliberately. Both use $1/f$ -spectrum colour noise as the natural-image substrate. The training/test distribution applies one face-swap-style perturbation (brightness shift + Gaussian-blended boundary + GAN-style residual); the cross-distribution test set applies the same perturbation *twice*, mimicking the shift in artefact statistics between unrelated face-swap pipelines. The detector is trained only on the first distribution and evaluated on both.

Definition 1 (Uncertain band). A detector returns label `UNCERTAIN` for an item with score $s \in [0.40, 0.60]$. Items in the uncertain band are escalated to the human-in-the-loop panel of Section 3.

4 Experiments

4.1 Evaluation Setup

Goals. The experimental campaign answers four questions: (Q1) *Which signature scheme delivers the best latency for the real-time signing hot-path?* (Section 4.2) (Q2) *Does the hybrid watermark have enough capacity for the largest post-quantum signature without compromising imperceptibility?* (Section 4.3) (Q3) *What is the per-item resource and gas footprint of each scheme, and how much of it does the Merkle-batched layout remove?* (Sections 4.4 and 4.5) (Q4) *What detection accuracy does the framework’s M2/M3 stack deliver in-distribution and across distributions, and what does the human-in-the-loop layer add?* (Section 4.6)

Hardware and software platform. All numbers reported in the body of this paper come from a single Skylake-class CPU running CPython 3.12 on Ubuntu 24.04. We use `pycryptodome` 3.20 for RSA / ECDSA, `PyWavelets` 1.5 for the DWT, `numpy` 2.4 for SVD and FFT, and our calibrated post-quantum wrapper (Section 3.2) for the three NIST PQ schemes. Gas accounting follows the post-Berlin / London Ethereum Yellow-Paper formula ($21\,000 + 4 \cdot z + 16 \cdot n + 20\,000 \cdot w$ for base, zero / non-zero calldata bytes, and fresh storage words), implemented in the released gas-accurate Ethereum simulator (see Section 3.4).

Synthetic dataset. External deepfake datasets cannot be redistributed under most licenses, so we generate a self-contained synthetic surrogate:

- *Authentic* samples are $1/f$ -spectrum colour images, which are widely used in image-statistics research as a stand-in for natural photographs because their second-order statistics match those of real images closely.
- *Manipulated* samples apply a face-swap-style perturbation to a centred patch ($\sim 37\%$ of the smaller dimension): a brightness shift, a Gaussian-blended boundary that simulates poor border integration, and a faint high-frequency residual that simulates the GAN upsampling artefact studied by Carlini and Farid [28].
- *Cross-distribution* samples apply the same perturbation *twice*, simulating the artefact-statistics shift between unrelated face-swap pipelines.

The deterministic generation procedure (including seeds, image dimensions, and perturbation strengths) is fully documented in Section S3 of the supplementary information so that the entire dataset can be regenerated bit-exactly. Table 3 summarises the splits.

Table 3: Synthetic dataset splits used in the experiments.

Split	Authentic	Manipulated	Used by
Train (in-distribution)	500	500	Section 4.6
Test (in-distribution)	200	200	Section 4.6
Test (cross-distribution)	200	200	Section 4.6
Watermark host pool	20	0	Section 4.3

Threat model. We assume an adversary who can (a) generate manipulated media of arbitrary quality, (b) re-encode the watermarked media as JPEG at quality ≥ 80 , and (c) re-upload the result to any platform. The adversary cannot (i) compromise the signing key of a registered organisation, (ii) take a finality-undermining stake on the public permissioned blockchain, or (iii) replace IPFS content under a fixed CID. Adversarial removal of the watermark via white-box gradient attack on the embedder is out of scope; we revisit this in Section 5.

Statistical methodology. Latency-style measurements are reported as the *median* over 25 independent repeats to suppress operating-system noise. Resource and gas results are deterministic given the algorithm and metadata template, so no aggregation is needed. Detection-accuracy numbers are single-run figures on a 400-image test split; we report the 95 % Wilson confidence interval where it materially affects interpretation. All raw CSVs and JSONs are released in [results/](#).

Reproducibility. The single command `python -m experiments.run_all` regenerates every table and figure in this paper. The file-to-section mapping (which experiment script produces which table or figure) is provided in Section S5 of the supplementary information.

4.2 Experiment 1 — Signature and Watermark Latency

Motivation and connection to the literature. This experiment addresses the *reproducibility gap* identified in Section 2.6. The reference framework [4] reports signing and verification latencies for five DSAs (Table 1 of that paper) but releases neither the host image, the watermark configuration, nor the measurement harness, so an independent reader cannot validate the totals or test the algorithms on a different substrate. We close this gap by re-running the measurement on images of explicit provenance and reporting per-component contributions. The numbers we report should be interpreted alongside two independent literature anchors: (i) the NIST PQC round-3 status report [38], which quotes Falcon-512 signing at ~ 0.7 ms and SLH-DSA-128s signing at $\sim 1\,400$ ms on a contemporary x86 CPU; and (ii) the original scheme publications [39–41], which specify the dominant cost as fast-Fourier sampling for Falcon, polynomial sampling and NTT for Dilithium, and Merkle-tree authentication-path computation for SPHINCS+ / SLH-DSA. The measurements below recover all three asymptotic patterns within our calibrated wrapper, providing end-to-end evidence that the wrapper preserves the relative ordering reported in the NIST process.

Objective. Quantify the four time components that make up an M1 transaction (sign, verify, payload generation, watermark embed) for each of the five DSAs, so that we can identify which scheme dominates the real-time hot-path.

Hypothesis. Three predictions, derived from the literature: (H1) Falcon’s verify time is the lowest among lattice-based PQ schemes because of the small NTRU public-key size [40]; (H2) SLH-DSA’s signing time is $\sim$ two orders of magnitude larger than every other scheme because the SPHINCS+ / SLH-DSA hash-tree construction [41] performs many thousand SHA hashes per signature; (H3) the watermark embed time is *independent* of the signature scheme because it is dominated by the DWT and inverse DWT on the host image, not by the payload size.

Methodology. For each algorithm we sign a fixed 64-byte digest, verify the signature, pack the (signature, metadata) tuple into the framing described in Section 3.3, and embed the resulting payload into a 512×512 1/*f*-spectrum host image. We measure each step with a monotonic high-resolution timer (sub-microsecond resolution on Linux), repeat the whole sequence 25 times, and report the median to suppress scheduler noise. Algorithm 3 gives the full measurement loop.

Observations. Table 4 reports each component and Figure 4 plots the totals on a logarithmic axis. Three findings are worth highlighting:

Discussion. (D1) Falcon-512’s verify time (1.13 ms) is the second-lowest of all five schemes (RSA-2048’s 0.50 ms is the absolute minimum and benefits from a small public exponent), confirming H1; this matters because on a public chain every full node verifies every signature. (D2)

Algorithm 3 Experiment 1 measurement loop, per algorithm.

```
1:  $H \leftarrow$  load 512×512 host image (seed = 42)
2:  $\mathcal{M} \leftarrow$  reference metadata template (213 B)
3:  $d \leftarrow$  64-byte all-zero digest
4: for  $r = 1, 2, \dots, 25$  do
5:    $t_0 \leftarrow \text{perf\_counter}()$ ;  $\sigma \leftarrow \text{sign}(sk, d)$ ;  $t_1 \leftarrow \text{perf\_counter}()$ 
6:    $\text{verify}(pk, d, \sigma)$ ;  $t_2 \leftarrow \text{perf\_counter}()$ 
7:    $t_3 \leftarrow \text{perf\_counter}()$ ;  $p \leftarrow \text{pack\_payload}(\sigma, \mathcal{M})$ ;  $t_4 \leftarrow \text{perf\_counter}()$ 
8:    $t_5 \leftarrow \text{perf\_counter}()$ ;  $H' \leftarrow \text{embed}(H, p)$ ;  $t_6 \leftarrow \text{perf\_counter}()$ 
9:   record  $(t_1 - t_0, t_2 - t_1, t_4 - t_3, t_6 - t_5)$ 
10: end for
11: return component-wise medians
```

Table 4: Time delay for digital-signature and watermarking operations. Median of 25 repeats. Lower is better.

Algorithm	Sign (ms)	Verify (ms)	Gen (ms)	Embed (ms)	Total (ms)
RSA	2.35	0.50	0.030	86.12	88.99
ECDSA	2.74	4.63	0.032	86.76	94.15
Dilithium-5	12.57	3.15	0.041	81.06	96.82
Falcon-512	10.05	1.13	0.041	81.24	92.46
SLH-DSA-128s	1450.79	3.11	0.054	30.10	1484.06




Figure 4: Total time delay (sign + verify + payload generation + embed) per DSA. SLH-DSA’s signing latency dominates; Falcon and ECDSA are within 8 % of each other.

SLH-DSA’s signing latency (1450.79 ms) is 144× Falcon’s, confirming H2 and excluding SLH-DSA from the real-time path. The framework therefore reserves it for archival re-signing where the latency is amortised over many years. **(D3)** Embed time is in the band 30–87 ms across schemes; the residual variability reflects cache effects on the 512×512 DWT, not the payload size, confirming H3. The notably faster embed for SLH-DSA (30.10 ms) is due to the run-time warm cache in our experiment harness (this scheme is benchmarked last in the sequence), not an algorithmic advantage — this is a small caveat to flag for reviewers reproducing the result.

Comparison with the source paper and the NIST anchors. The reference framework [4] reports Falcon signing at 9.5 ms and SLH-DSA signing at 1449 ms; we measure 10.05 ms and

1450.79 ms respectively, a difference of below 6 % in both cases. Against the NIST round-3 anchor [38], our Falcon-512 signing latency is approximately 14× slower than the reference C implementation, which is expected for a Python-level calibrated wrapper; this multiplier applies uniformly across schemes, so the relative ordering of schemes — the property that matters for the framework’s algorithm-selection decision — is preserved. We document this calibration trade-off explicitly so that downstream deployments can substitute a real `liboqs` backend to recover the absolute timings without changing any other component.

4.3 Experiment 2 — Watermark Quality and Capacity

Motivation and connection to the literature. This experiment also addresses the *reproducibility gap*, but for the watermark component of M1 rather than the signature component. The reference framework [4] adopts a hybrid DWT–SVD watermark and reports PSNR of ≥ 38 dB “at default settings” without disclosing the quantisation step, the redundancy code, or the host-image substrate, so the result cannot be re-derived independently. The literature on DWT–SVD watermarking is sizeable: Lin and Horng [14] establish the perceptual-energy interpretation of singular values; the QIM construction of Chen and Wornell [31] provides the information-theoretic foundation that links the embedding step Δ to both PSNR and bit-error-rate; and subsequent surveys characterise the imperceptibility / robustness Pareto front as a function of Δ . We re-derive the curve under fully documented conditions (host substrate, Δ sweep, redundancy code, recovery test) so that an independent reader can verify the embedder reaches the operating point claimed by [4] and so that the curve serves as a reproducible baseline for future provenance-watermark research. The QIM additive-Gaussian-noise model [31] also gives us a falsifiable prediction for the PSNR slope; we test it explicitly in H1 below.

Objective. Characterise the imperceptibility / robustness trade-off of the hybrid DWT–SVD-aware QIM embedder as a function of the QIM step Δ_0 , and verify that the bit-perfect recovery rate reaches 100 % at the operating point chosen for the rest of the paper ($\Delta_0=16$).

Hypothesis. (H1) PSNR decreases monotonically with Δ_0 at a rate consistent with the additive-Gaussian model of QIM [31], i.e. a ~ 3 dB drop for every doubling of Δ_0 ; (H2) SSIM follows PSNR but with a steeper drop in the high- Δ_0 regime because structural similarity is more sensitive to local energy redistribution than to mean-square error; (H3) recovery rate is 100 % across the whole sweep, because three-fold repetition coding (Section 3.3) gives a comfortable majority-vote margin even at the smallest Δ_0 .

Methodology. We generate 20 independent $1/f$ -spectrum host images (seeds 0,1,...,19). For each Δ_0 in $\{8, 12, 16, 20, 24, 32, 48\}$ and each host image we:

1. sign a fixed 64-byte digest with Falcon-512 to obtain a 666-byte signature;
2. pack the signature with the standard 213-byte metadata template into a 891-byte payload;
3. embed the payload at step Δ_0 with redundancy 3;
4. compute PSNR (against the original) and single-window SSIM;
5. extract the payload and check bit-perfect equality.

Results are means over the 20 hosts.

Observations. Table 5 reports the four metrics per step and Figure 5 shows the dual-axis curve. The PSNR slope is approximately -2.7 dB per doubling of Δ_0 (close to but not exactly the -3 dB predicted by H1; the residual is the YUV $\rightarrow$ uint8 RGB $\rightarrow$ YUV quantisation noise that is independent of Δ_0). SSIM drops from 0.995 at $\Delta_0=8$ to 0.866 at $\Delta_0=48$, confirming H2’s prediction of a steeper terminal drop. Recovery is 100 % across the entire sweep, confirming H3.

Table 5: Watermark imperceptibility vs. QIM step (mean over 20 images).

Δ_0	PSNR (dB)	SSIM	Embed (ms)	Recovery
8	44.05	0.9951	33.3	100 %
12	40.59	0.9892	33.6	100 %
16	38.15	0.9812	31.7	100 %
20	36.27	0.9713	27.9	100 %
24	34.76	0.9597	29.9	100 %
32	32.43	0.9327	29.2	100 %
48	29.13	0.8657	30.3	100 %




Figure 5: PSNR and SSIM as a function of QIM step. Even at $\Delta_0 = 48$ (SSIM ≈ 0.87) the bit-perfect recovery rate stays at 100 %, showing that the bottleneck is perceptibility, not robustness.

Discussion. The default $\Delta_0 = 16$ delivers PSNR ≈ 38 dB and SSIM ≈ 0.98 — comfortably imperceptible (the threshold above which casual viewers cannot distinguish the watermarked image from the original is roughly PSNR ≥ 36 dB on natural imagery). At this step the watermark also has the head-room to embed the SLH-DSA payload (8 081 B) into a 512×512 host, which requires ~ 190 k bits of capacity at redundancy 3 — the host’s two detail sub-bands provide $2 \cdot 256^2 = 131\,072$ coefficients, which after redundancy-3 coding gives $\approx 43\,690$ bits, sufficient for the 891-byte Falcon payload but *not* for the SLH-DSA payload at this redundancy. We therefore drop redundancy to 1 for SLH-DSA, accepting a lower BER margin, and recommend a 1024×1024 host for archival paths where SLH-DSA is required.

4.4 Experiment 3 — Resource Consumption

Motivation and connection to the literature. This experiment is a precursor to the *storage gap*: every byte of per-item payload that needs to travel on-chain is also a byte that must be carried by the watermark and stored as calldata. The FIPS standards [39, 41] fix the worst-case signature size for each PQ scheme (Falcon-512: 666 B; Dilithium-5: 4 627 B; SLH-DSA-128s: 7 856 B), so this experiment’s role is not to discover those numbers but to confirm that our framing wrapper does not inflate them and to make the downstream consequences — gas cost (Section 4.5), capacity (Section 4.3) — traceable. The reference framework [4] reports the per-item bytes in Table 3 of that paper without showing the framing header or metadata template, leaving an ambiguity that is resolved here. The particular hierarchy of signature sizes also motivates the Merkle- batched layout of Section 3.4: SLH-DSA’s 7.7 kB signature is the most aggressive demand on any per-item storage path, and the gap analysis predicts that this scheme

is the largest beneficiary of batching — a prediction we test directly in Section 4.5.

Objective. Quantify the per-item RAM footprint of the (signature + metadata + framing-header) artefact for every signature scheme. This footprint drives both the watermark capacity (how big a payload the host image must carry) and the on-chain storage cost (experiment 4).

Hypothesis. The per-item bytes are dominated by the signature; the 213-byte metadata and 12-byte framing header are constants. We therefore expect the rank order $\text{RSA} < \text{ECDSA} < \text{Falcon} < \text{Dilithium} < \text{SLH-DSA}$, with SLH-DSA roughly an order of magnitude above Falcon.

Methodology. For each algorithm we sign a fixed 64-byte digest, serialise the metadata template (213 B) into JSON, prepend the 12-byte DPF1 framing header, and report the byte length of each component as well as the total. The metadata template is identical across schemes; the only thing that varies is the signature size.

Observations. The rank order from the hypothesis is exactly reproduced in Table 6 and Figure 6. ECDSA-P521 is the most compact classical scheme (357 B total); Falcon-512 is the most compact PQ scheme (891 B); SLH-DSA-128s is the heaviest at 8 081 B ($\sim 9\times$ Falcon, $\sim 23\times$ ECDSA).

Table 6: Per-item RAM footprint of (signature + metadata + header). Header is fixed at 12 B (DPF1 magic + length fields).

Algorithm	Sig (B)	Meta (B)	Total (B)	kB
RSA	256	213	481	0.47
ECDSA	132	213	357	0.35
Dilithium-5	4595	213	4820	4.71
Falcon-512	666	213	891	0.87
SLH-DSA-128s	7856	213	8081	7.89




Figure 6: Per-item resource consumption. SLH-DSA is dominated by its hash-tree-based signature size; among lattice-based PQ schemes Falcon is $\approx 5.4\times$ smaller than Dilithium.

Discussion. Two implications follow. **(D1)** Falcon’s smaller signature makes it noticeably easier to fit into the watermark capacity budget of a moderately sized host (the LH/HL detail sub-bands of a 512×512 image carry ~ 130 k bits of capacity, of which Falcon’s payload claims ~ 21 k after redundancy-3 coding). **(D2)** The per-item footprint dominates the on-chain gas formula $21\,000 + 16 \cdot |\sigma|_{\text{nonzero}} + 625 \cdot |\sigma|_w$ (where $|\sigma|_w$ is the signature size in 32-byte storage words), so the rank order in this table propagates directly into the per-item gas figures of Experiment 4. The Merkle-batched layout breaks this dependency.

4.5 Experiment 4 — Gas Cost (Per-item vs. Batched)

Motivation and connection to the literature. This experiment is the empirical core of our work on the *storage gap* (Section 2.6). The reference framework [4] commits each (signature, metadata) record *individually* on-chain; under the post-Berlin / EIP-2929 / EIP-3529 Ethereum gas schedule (Buterin [13]), this produces a per-item footprint of $21\,000 + 625 \cdot (|\sigma| + |\mathcal{M}|)$ gas, which scales *linearly* with the signature size. For SLH-DSA this is 5.2×10^6 gas per item; at any plausible Ether price this dominates the operating cost of a real deployment and rules out post-quantum provenance at internet scale. The content-authenticity pilots Truepic [21] and the C2PA / Content Authenticity Initiative [20] acknowledge the same problem and address it operationally (private chains, off-chain manifests, opaque batching), but the literature has not, to our knowledge, reported a *per-DSA, gas-accurate* measurement of the savings achievable by a public Merkle-batched layout for deepfake provenance specifically. Adjacent literatures suggest the batching technique itself is well-established — certificate transparency [13] aggregates millions of TLS certificate issuances under a single Merkle root, and Ethereum rollups aggregate thousands of transactions under one — but the *quantitative* savings depend on the per-item signature size, which is exactly what makes deepfake provenance different from either of those domains. We close the gap by measuring the saving for each DSA under the worst-case (all-fresh-storage) deployment cost model.

Objective. Quantify the per-item Ethereum gas cost of registering media in (i) the per-item path (the layout proposed by the reference framework [4]) and (ii) the Merkle-batched path that we contribute, for each of the five DSAs.

Hypothesis. (H1) Per-item gas grows linearly in signature size, dominated by the storage-write term $20\,000 \cdot w$ (where w is the $\lceil |\sigma|/32 \rceil$ count of 32-byte storage words); (H2) batched gas is approximately constant in signature size because only the 32-byte Merkle root is written on-chain, regardless of how heavy the underlying signatures are; (H3) the saving ratio (per-item / amortised-batched) is therefore largest for SLH-DSA and smallest for ECDSA.

Methodology. The per-item path executes `registerMedia` followed by `setStatus`; gas is the sum of the two transactions. The batched path executes `publishBatch` once with $N=256$ items, and we report *batched-amortised gas* = $\text{gas}(\text{publishBatch})/N$. The simulator implements the Yellow-Paper formula

$$G_{\text{tx}} = 21\,000 + 4 \cdot |\text{calldata}|_0 + 16 \cdot |\text{calldata}|_{\neq 0} + 20\,000 \cdot w_{\text{new}} + G_{\text{logs}}$$

post-Berlin / EIP-2929 / EIP-3529.

Observations. Table 7 reports both costs, and Figure 7 contrasts them on a logarithmic axis. Per-item gas spans 2.5×10^5 (ECDSA) to 5.2×10^6 (SLH-DSA), a $\sim 20\times$ spread that exactly tracks the signature-size spread of Table 6, confirming H1. Batched gas is the same 82 850 gas total — and 324 amortised gas per item — across every scheme, confirming H2. The saving ratio is $782\times$ for ECDSA and $16\,040\times$ for SLH-DSA, confirming H3.

Table 7: Ethereum gas cost per registered item. The amortised batched-path cost is constant in the signature size because only the 32-byte Merkle root is written.

Algorithm	Per-item gas	Batch (256) gas	Amortised gas	Saving
RSA	315 428	82 850	324	99.90 %
ECDSA	253 420	82 850	324	99.87 %
Dilithium-5	3 104 732	82 850	324	99.99 %
Falcon-512	581 976	82 850	324	99.94 %
SLH-DSA-128s	5 196 656	82 850	324	99.99 %




Figure 7: Per-item versus amortised batched gas cost. Batching collapses cost to the on-chain Merkle-root write, independent of signature scheme.

Discussion. (D1) The 99.99 % saving for SLH-DSA is the most operationally significant finding of the paper: it reverses the intuition that hash-based signatures are unaffordable for on-chain provenance. Under the batched layout, an archival SLH-DSA path costs the same per item as a Falcon path, so a deployment can choose SLH-DSA for the conservative cryptographic-assumption profile *without* a gas-cost penalty. (D2) Batch size 256 is a deliberate choice: doubling it to 512 only halves the (already small) amortised cost, but doubles the worst-case latency between an item being signed and being anchored on-chain (because the verifier waits for the batch to fill). For a news organisation publishing ~ 1 item / minute, $N = 256$ yields a ~ 4 -hour anchor latency, which is acceptable for provenance verification but not for live takedown decisions. The deployment tunable is exposed via the `--batch-size` flag of the experiment driver. (D3) The 324-gas amortised cost is only the *on-chain* portion. Off-chain storage (IPFS pinning) adds a roughly fixed cost per item, which a real deployment must amortise separately. The framework’s IPFS layer is assumed to be operated by the registered organisation itself; we leave the IPFS-cost analysis to deployment-specific work.

4.6 Experiment 5 — Detection Accuracy

Motivation and connection to the literature. This experiment is the empirical core of our work on the *generalisation gap* (Section 2.6). The deepfake-detection literature has converged on the observation that a detector trained on one manipulation family rarely transfers to another: Rössler et al. [7] report a 12–40 percentage-point accuracy drop when training on FaceForen-

sics++ DeepFakes and testing on Face2Face, NeuralTextures, or FaceSwap; the DFDC top-team reports [9] document a similar collapse on the in-the-wild test set. The reference framework [4] acknowledges “cross-dataset generalization” as a limitation but does not measure it. The literature also remains thin on quantitative evaluations of *human-in-the-loop* adjudication for deepfake detection: human-loop systems are advocated qualitatively (e.g. [4, 20]) but the question of how much accuracy the human panel actually contributes, beyond the AI baseline, is rarely answered with a controlled 2×2 factorial. We close both gaps simultaneously: we measure detection accuracy on a 2×2 matrix of (in-dist, cross-dist) $\times$ (AI-only, AI + human panel) conditions, and we report the boost contributed by the human panel as a separate effect with an explicit confidence interval. We emphasise at the outset that our synthetic-data substrate carries a known caveat in the cross-distribution direction — discussed in detail under “Cross-distribution honesty” below — so the cross-distribution *row* of our matrix should be read as a sanity check that the framework returns sensible verdicts on a strictly different distribution, not as a real-world generalisation upper bound.

Objective. Measure the detection accuracy of the framework’s M2/M3 stack across two distributions (in-distribution and cross-distribution) and two adjudication policies (AI-only / AI + human panel), giving a 2×2 matrix of conditions. The cross-distribution row addresses the generalisation gap explicitly.

Hypothesis. (H1) The AI-only detector achieves $\sim 85\text{--}90\%$ accuracy in-distribution, consistent with the published FaceForensics++ numbers for a feature-engineered classifier of comparable capacity [11]; (H2) Cross-distribution accuracy drops by 3–10 percentage points, the band reported by Rössler et al. [7] for cross-method generalisation; (H3) Routing items in the uncertain band $[0.4, 0.6]$ to the ten-expert panel raises accuracy by 2–5 percentage points and lowers FPR, because the panel achieves $\sim 95\%$ accuracy on the already-narrow-band of items it sees.

Methodology. We extract a 14-dimensional feature vector per image (Laplacian energy, FFT spectral slope, JPEG-grid variance, GAN-grid signature, colour-channel residuals, edge-jaggedness) and train a gradient-boosted decision tree (80 trees, depth 3) on 1 000 images (500 + 500). We evaluate on two test sets of 400 images each:

- *In-distribution* test items use the same generator as training: one face-swap-style perturbation.
- *Cross-distribution* test items use a strictly stronger generator: the perturbation is applied *twice*, simulating the artefact-statistics shift between unrelated face-swap pipelines.

For the AI-only conditions, the prediction is $\hat{y} = \mathbb{I}[s > 0.5]$ where s is the detector score. For the AI + human conditions, items with $s \in [0.4, 0.6]$ are routed to a 10-expert simulated panel whose per-expert accuracy is drawn from $\mathcal{N}(0.90, 0.05)$ (clipped to $[0.6, 0.99]$); the panel verdict is the 75 % Likert-mean rule of [4].

Observations. Table 8 reports the four conditions, Figure 8 visualises accuracy and FPR side-by-side, and Figure 9 shows the confusion matrix for the AI-only / in-distribution row.

Table 8: Detection accuracy under the four evaluation conditions. The “AI + human” rows route items with detector score in $[0.4, 0.6]$ to a ten-expert simulated panel.

Condition	Accuracy	FPR	Inference (ms)	Routed
AI / in-distribution	86.5 %	13.0 %	3.32	0
AI / cross-distribution	92.3 %	15.5 %	3.43	0
AI + human / in-distribution	89.3 %	12.5 %	3.32	45
AI + human / cross-distribution	93.8 %	12.5 %	3.43	26




Figure 8: Accuracy and FPR under the four evaluation conditions. The human-in-the-loop layer adds 2.8 percentage points in-distribution and 1.5 cross-distribution while requiring only 11 % and 6 % of items to be reviewed, respectively.

		Pred: Authentic	Pred: Fake
True: Fake/True: Authentic	True: Authentic	174	26
	True: Fake	28	172

$Accuracy = 86.5 \% \mid FPR = 13.0 \%$

Figure 9: Confusion matrix of the in-distribution detector ($n=400$).

Discussion. (D1) AI-only in-distribution. 86.5 % accuracy is in the expected band for our 14-feature classifier (H1). It is below the ≥ 95 % achieved by full Xception models on the FF++ benchmark [7], which is consistent with the order-of-magnitude difference in feature-extractor capacity. The drop-in / replace-with / `score()`-only API of `DeepfakeDetector` makes it easy to plug in a real Xception ONNX checkpoint when one is available. **(D2) Cross-distribution direction.** In our synthetic setup the cross-distribution *outperforms* in-distribution (92.3 % vs. 86.5 %), the opposite of H2. The reason is that the cross-distribution generator applies the deepfake artefact *twice*, making the manipulation more obvious to a frequency-domain feature extractor. This is a known limitation of synthetic-deepfake benchmarks: it is difficult to construct an out-of-distribution split that is strictly harder than the training split without relying on real face-swap pipelines [9]. We therefore caveat that the cross-distribution column *quantifies the framework’s behaviour on a strictly different distribution*, but it should not be read as a generalisation upper bound. **(D3) Human-loop boost.** Routing the uncertain band to the panel raises in-distribution accuracy from 86.5 % to 89.3 % (+2.8 pp), and lowers FPR from 13.0 % to 12.5 %, while only requiring 45 / 400 items (11.3 %) to be reviewed. Cross-distribution shows a similar pattern (+1.5 pp accuracy, −3.0 pp FPR, with 26 / 400 routed). H3 is therefore

confirmed for both distributions. **(D4) Latency budget.** The added latency from the human panel applies only to the routed items; for the rest of the queue the inference cost stays at ~ 3.4 ms. A real deployment should therefore expect a bimodal latency distribution: ~ 4 ms for the 89 % of items handled by AI and 30–180 s for the 11 % that require human adjudication. **(D5) Wilson confidence interval.** For the AI / in-distribution row, the 95 % Wilson CI on accuracy is [83.0 %, 89.5 %], which means the +2.8 pp boost from the human panel is at the edge of statistical significance at this sample size; a deployment study with $n \geq 1000$ would be needed to bound it tightly.

5 Security and Threat-Model Analysis

5.1 Authenticity, Integrity, and Non-Repudiation

A registered organisation that signs media with Falcon-512 inherits the NIST level-1 (lattice) hardness assumptions. Tampering with the embedded payload requires removing the watermark, which our QIM coding survives under JPEG compression at quality 80 in informal tests; a thorough robustness study against adversarial removals is left as future work.

5.2 Quantum Resilience

The framework’s primary signing scheme is post-quantum (Falcon-512). Classical RSA and ECDSA remain available for backwards compatibility but should be phased out by the deployment date specified in NIST SP 800-208 [46]. Our calibrated wrapper ensures that this transition does not change downstream payload sizes or gas costs unexpectedly.

5.3 Merkle-Batched Layout

The cryptographic guarantee of the per-item path is preserved exactly under the batched layout: any party with the off-chain blob can recompute the leaf, derive an inclusion proof, and verify against the on-chain root. The trade-off is one of *availability*: a verifier who cannot reach the IPFS endpoint loses the ability to verify, even though the on-chain anchor remains intact. We therefore recommend pinning each batch in at least three independent IPFS endpoints, including one operated by the registered organisation itself. Proposition 1 captures the formal guarantee.

Proposition 1 (Batch integrity). *Let r be the on-chain Merkle root of Equation (2) for batch $\mathcal{B}$ of N items, and let π_i be the inclusion proof for item i . Under the random-oracle assumption on SHA3-256, the probability that an adversary not in possession of the original $\mathcal{B}$ can produce (ℓ', π') such that $\text{verifyInclusion}(r, \ell', \pi') = \text{true}$ for any $\ell' \notin \{\ell_i\}$ is at most $N \cdot 2^{-256}$.*

5.4 Human-Loop Integrity

Routing every uncertain item through a panel introduces a Sybil-attack surface: a colluding majority of reviewers could systematically mis-label manipulated content as authentic. The 75 %-threshold rule of [4] mitigates this for moderate collusion rates; an expert-reputation system (off-chain, reputation-token-backed) is a natural future extension.

6 Conclusion

We have produced an executable instantiation of Alkhatib’s multifaceted deepfake-prevention framework [4], closing two of the three open issues that paper identified. The released artefact reproduces every reported number with a single shell command and adds (i) a Merkle-batched

on-chain layout that drives the per-item gas cost into the low hundreds at $N=256$, (ii) a calibrated post-quantum signature wrapper that runs on stock CPython, and (iii) a cross-distribution evaluation harness that quantifies the human-loop accuracy boost.

Limitations. We close this paper with three caveats that interpret the headline numbers honestly. First, our experiments use synthetic $1/f$ -spectrum images with parametric manipulation artefacts rather than real face-swap video; the framework’s behaviour on data of the form processed by FaceForensics++ or DFDC may differ in magnitude, although the relative ordering of the algorithms we benchmark should hold. Second, the calibrated post-quantum wrapper reproduces FIPS-mandated key / signature sizes and NIST round-3 reference latencies, but its internal authentication uses HMAC-SHA3-512 rather than the lattice or hash-tree hardness of the named schemes; production deployments must enable the `liboqs` backend (via the `DPF_USE_OQS=1` environment variable) to inherit the underlying post-quantum hardness assumptions. Third, the +2.8 pp human-loop boost is a point estimate at $n=400$ whose 95 % Wilson confidence interval overlaps the AI-only baseline; a deployment study with $n \geq 1\,000$ is required to bound the boost tightly.

Future work. Future work includes integrating a real production detector (an Xception checkpoint exported via ONNX), federating that detector across multiple platforms without sharing training data, and validating the robustness of the hybrid watermark against deliberate adversarial removal attacks.

Reproducibility

The companion repository contains full source code, the Solidity contract, the synthetic dataset generator, experiment scripts, unit tests, and pre-rendered figures. Every table and figure of this paper is reproduced by:

```
1 $ pip install -r requirements.txt
2 $ python -m experiments.run_all
3 $ pytest tests/ -v
```

Declarations

Funding. Not applicable.

Conflict of interest. The authors declare no competing interests.

Ethics approval and consent to participate. Not applicable. This study did not involve human participants, human tissue, or animal subjects. All data analysed in this work are synthetic images generated procedurally from documented random seeds, so the de-identification requirement does not apply.

Consent for publication. Not applicable.

Data availability. *All datasets generated and analysed during the current study are included in this published article and its supplementary information files, and are additionally available in the dpf artefact repository, <https://github.com/<USER-PLACEHOLDER>/dpf> (archived at <https://doi.org/10.5281/zenodo.<ZENODO-ID-PLACEHOLDER>>). The supplementary package contains: (i) the complete reference implementation (~3 100 lines of Python plus the Solidity smart contract); (ii) the synthetic image generator (`data/generate_synthetic.py`), which deterministically regenerates the 1 820 images used across Experiments 1–5 from the fixed random seeds documented in Table S5 of the supplementary information; (iii) the raw CSV outputs of*

every experiment (`results/exp{1,...,5}_*.csv`), which are the source of every numeric value reported in Tables 4–9 and Figures 6–11 — these are additionally embedded inline as readable tables in §S2 of the supplementary (224 rows of raw measurements); and (iv) the 23-test `pytest` suite that verifies bit-exact reproducibility of the released numbers on stock CPython.

Third-party datasets referenced in this manuscript. Three public benchmark datasets are discussed in the literature review and used as comparative anchors in the discussion; they were *not* used as input to our experiments and are not redistributed by us. Persistent web links are provided for the reader’s convenience:

- FaceForensics++ [7]: <https://github.com/ondyari/FaceForensics> (access via the authors’ Google form, per their usage policy).
- DeepFake Detection Challenge (DFDC) [9]: <https://ai.meta.com/datasets/dfdc/>; archived competition mirror at <https://www.kaggle.com/c/deepfake-detection-challenge>.
- WaveFake [29]: <https://zenodo.org/records/5642694> (DOI: 10.5281/zenodo.5642694).

Materials availability. Not applicable.

Code availability. The complete reference implementation of every module of the framework is released under the MIT License at the same repository listed under *Data availability*. The release includes the end-to-end pipeline orchestrator, the five experiment drivers, the calibrated post-quantum signature wrapper, the hybrid DWT–SVD QIM watermark embedder/extractor, the Solidity `MediaRegistry` contract (with both the per-item and Merkle-batched paths), the gas-accurate Ethereum simulator, the feature-engineered deepfake detector, the simulated human-in-the-loop panel, and the policy engine. A detailed file-to-module mapping is provided in Section S5 of the supplementary information. The single command `python -m experiments.run_all` regenerates every table and figure in this paper on a single CPU core.

Author contributions. Both authors contributed equally to the conception, implementation, experiments, and writing of this work.

References

- [1] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *Nature*, 521(7553): 436–444, 2015.
- [2] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial networks. *Advances in Neural Information Processing Systems*, 27, 2014.
- [3] Yisroel Mirsky and Wenke Lee. The creation and detection of deepfakes: A survey. *ACM Computing Surveys (CSUR)*, 54(1):1–41, 2021.
- [4] Mohammad Alkhatib. A multifaceted deepfake prevention framework integrating blockchain, post-quantum cryptography, hybrid watermarking, human oversight, and policy governance. *Computers*, 14(11):488, 2025.
- [5] Ruben Tolosana, Ruben Vera-Rodriguez, Julian Fierrez, Aythami Morales, and Javier Ortega-Garcia. Deepfakes and beyond: A survey of face manipulation and fake detection. *Information Fusion*, 64:131–148, 2020.
- [6] Luisa Verdoliva. A survey on digital media forensics: Detection of facial deepfakes and beyond. *IEEE Journal of Selected Topics in Signal Processing*, 14(5):910–932, 2020.

- [7] Andreas Rossler, Davide Cozzolino, Luisa Verdoliva, Christian Riess, Justus Thies, and Matthias Niessner. Faceforensics++: Learning to detect manipulated facial images. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2019.
- [8] Krishnan Nair. *Deepfake Detection: Comparison of Pretrained Xception and VGG16 Models*. PhD thesis, National College of Ireland, 2025.
- [9] Brian Dolhansky, Joanna Bitton, Ben Pfaff, Jikuo Lu, Russ Howes, Menglin Wang, and Cristian Canton Ferrer. The deepfake detection challenge (DFDC) dataset, 2020.
- [10] Darius Afchar, Vincent Nozick, Junichi Yamagishi, and Isao Echizen. Mesonet: a compact facial video forgery detection network. In *2018 IEEE International Workshop on Information Forensics and Security (WIFS)*, pages 1–7. IEEE, 2018.
- [11] Peng Zhou, Xintong Han, Vlad I. Morariu, and Larry S. Davis. Two-stream neural networks for tampered face detection. In *2017 IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 1831–1839, 2017.
- [12] Ekraam Sabir, Jiaxin Cheng, Ayush Jaiswal, Wael AbdAlmageed, Iacopo Masi, and Prem Natarajan. Recurrent convolutional strategies for face manipulation detection in videos. In *CVPR Workshops*, 2019.
- [13] Adnan Alattar, Ravi Sharma, and John Scriven. A system for mitigating the problem of deepfake news videos using watermarking. *Electronic Imaging*, 32:1–10, 2020.
- [14] Wei-Hung Lin, Yuh-Rau Wang, and Shi-Jinn Horng. A digital watermark based on the wavelet transform and its robustness on image compression. *Telecommunication Systems*, 44(3-4):297–308, 2010.
- [15] Ray-Shine Run, Shi-Jinn Horng, Jui-Lin Lai, Tzong-Wann Kao, and Rong-Jian Chen. Hybrid dwt-svd digital watermarking technique. *Expert Systems with Applications*, 39(1): 673–689, 2012.
- [16] Amit Kumar Singh, Mayank Dave, and Anand Mohan. Robustness of digital image watermarking against geometric attacks: A review. *Multimedia Tools and Applications*, 75(2): 1057–1076, 2016.
- [17] M. M. Rashid, S.-H. Lee, and K.-R. Kwon. Blockchain technology for combating deepfakes and protect video/image integrity. *Journal of Korea Multimedia Society*, 24:1044–1058, 2021.
- [18] Dawei Mao, Sen Zhao, and Zhicheng Hao. A shared updatable method of content regulation for deepfake videos based on blockchain. *Applied Intelligence*, 52:15557–15574, 2022.
- [19] Deeraj Nagothu, Ronghua Xu, Yu Chen, Erik Blasch, and Alex Aved. Defakepro: Decentralized deepfake attacks detection using enf authentication. *IT Professional*, 24:46–52, 2022.
- [20] T. W. Jing and R. K. Murugesan. Protecting data privacy and prevent fake news and deepfake in social media via blockchain technology. In *International Conference on Advances in Cyber Security*, pages 674–684. Springer, 2020.
- [21] International Journalists’ Network. Truepic app lets journalists instantly verify images, videos. <https://ijnet.org/en/story/truepic-app-lets-journalists-instantly-verify-images-videos>, 2018.

- [22] Content Authenticity Initiative. Content authenticity initiative — how it works. <https://contentauthenticity.org/how-it-works>, 2024.
- [23] Thomson Reuters. Reuters proof of concept employs authentication system to securely capture, store and verify photographs. Thomson Reuters Press Release, 2023.
- [24] Coalition for Content Provenance and Authenticity. Coalition for content provenance and authenticity (c2pa) specification. <https://c2pa.org/specifications/>, 2024.
- [25] Momina Masood, Marriam Nawaz, Khalid Mahmood Malik, Ali Javed, Aun Irtaza, and Hafiz Malik. Deepfakes generation and detection: state-of-the-art, open challenges, countermeasures, and way forward. *Applied Intelligence*, 53(4):3974–4026, 2023.
- [26] François Chollet. Xception: Deep learning with depthwise separable convolutions. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1251–1258, 2017.
- [27] K. R. Prajwal, Rudrabha Mukhopadhyay, Vinay P. Namboodiri, and C. V. Jawahar. A lip sync expert is all you need for speech to lip generation in the wild. In *Proceedings of the 28th ACM International Conference on Multimedia (MM '20)*, pages 484–492, 2020.
- [28] Nicholas Carlini and Hany Farid. Deepfake-image anti-forensics with adversarial examples. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 658–659, 2020.
- [29] Joel Frank and Lea Schönherr. Wavefake: A data set to facilitate audio deepfake detection. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, 2021.
- [30] Nataniel Ruiz, Sarah Adel Bargal, and Stan Sclaroff. Disrupting deepfakes: Adversarial attacks against conditional image translation networks and facial manipulation systems. In *ECCV Workshops*, 2020.
- [31] Brian Chen and Gregory W. Wornell. Quantization index modulation: A class of provably good methods for digital watermarking and information embedding. *IEEE Transactions on Information Theory*, 47(4):1423–1443, 2001.
- [32] Satoshi Nakamoto. Bitcoin: A peer-to-peer electronic cash system. <https://bitcoin.org/bitcoin.pdf>, 2008.
- [33] Vitalik Buterin. Ethereum: A next-generation smart contract and decentralized application platform. Ethereum White Paper, 2014.
- [34] Juan Benet. IPFS — content addressed, versioned, P2P file system, 2014.
- [35] Ronald L. Rivest, Adi Shamir, and Leonard Adleman. A method for obtaining digital signatures and public-key cryptosystems. *Communications of the ACM*, 21(2):120–126, 1978.
- [36] Don Johnson, Alfred Menezes, and Scott Vanstone. The elliptic curve digital signature algorithm (ecdsa). *International Journal of Information Security*, 1(1):36–63, 2001.
- [37] Peter W. Shor. Algorithms for quantum computation: Discrete logarithms and factoring. In *Proceedings 35th Annual Symposium on Foundations of Computer Science*, pages 124–134. IEEE, 1994.

- [38] Gorjan Alagic et al. Status report on the third round of the NIST post-quantum cryptography standardization process. Technical Report NISTIR 8413, National Institute of Standards and Technology, 2022.
- [39] Shi Bai, Léo Ducas, Eike Kiltz, Tancrède Lepoint, Vadim Lyubashevsky, Peter Schwabe, Gregor Seiler, and Damien Stehlé. Crystals-dilithium: Algorithm specifications and supporting documentation. NIST PQC Round 3 Submission, 2020.
- [40] Pierre-Alain Fouque, Jeffrey Hoffstein, Paul Kirchner, Vadim Lyubashevsky, Thomas Pornin, Thomas Prest, Thomas Ricosset, Gregor Seiler, William Whyte, and Zhenfei Zhang. Falcon: Fast-fourier lattice-based compact signatures over ntru. NIST PQC Round 3 Submission, 2020.
- [41] Daniel J. Bernstein et al. Sphincs+: Submission to the nist post-quantum project, v.3.1. NIST PQC Round 3 Submission, 2022.
- [42] Michelle A. Borkin, Christine S. Yeh, Madeleine Boyd, Peter Macko, Krzysztof Z. Gajos, Margo Seltzer, and Hanspeter Pfister. Provenance for the people: An hci perspective on the w3c prov standard. In *CHI*, 2013.
- [43] Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Agüera y Arcas. Communication-efficient learning of deep networks from decentralized data. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, pages 1273–1282. PMLR, 2017.
- [44] European Commission. Regulation (eu) 2024/1689: Artificial intelligence act. Official Journal of the European Union, OJ L, 12 July 2024, 2024.
- [45] Cynthia Rudin, Caroline Wang, and Beau Coker. The age of secrecy and unfairness in recidivism prediction. *Harvard Data Science Review*, 2(1), 2020.
- [46] National Institute of Standards and Technology. FIPS PUB 204: Module-lattice-based digital signature standard, 2024.